

```
# Boston Housing Price Prediction - ShadowFox Internship Task
# Author: Vedant

# =====
# 1. Import Required Libraries
# =====
from sklearn.datasets import fetch_california_housing
import pandas as pd

# Load California Housing dataset
housing = fetch_california_housing(as_frame=True)

# Create DataFrame
df = housing.frame
df.rename(columns={"MedHouseVal": "MEDV"}, inplace=True) # rename target column for consistency

# Save to CSV (optional)
df.to_csv("BostonHousing.csv", index=False)

print("Dataset shape:", df.shape)
print(df.head())

import pandas as pd
import numpy as np
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
```

```
# =====
# 2. Load Dataset
# =====
# NOTE: Place your BostonHousing.csv in the same folder as this notebook
df = pd.read_csv("BostonHousing.csv")
print("Dataset shape:", df.shape)
print(df.head())

# =====
# 3. Split Features and Target
# =====
X = df.drop("MEDV", axis=1) # Features
y = df["MEDV"]             # Target variable (House Price)

# =====
# 4. Train-Test Split
# =====
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
print("Train size:", X_train.shape, "Test size:", X_test.shape)

# =====
# 5. Data Normalization
# =====
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```

# 6. Build Regression Model (Neural Network)
# =====
model = tf.keras.Sequential([
    tf.keras.layers.Dense(64, activation="relu"),
    tf.keras.layers.Dense(32, activation="relu"),
    tf.keras.layers.Dense(1) # single output for regression
])

model.compile(optimizer="adam", loss="mse", metrics=["mae"])

# =====
# 7. Train Model
# =====
history = model.fit(X_train, y_train, epochs=100, validation_data=(X_test, y_test), verbose=0)

# =====
# 8. Evaluate Model
# =====
loss, mae = model.evaluate(X_test, y_test, verbose=0)
print(f"Test Mean Absolute Error: {mae:.2f}")

# =====|
# 9. Visualize Training
# =====
plt.plot(history.history['mae'], label='Train MAE')
plt.plot(history.history['val_mae'], label='Validation MAE')
plt.xlabel("Epochs")
plt.ylabel("Mean Absolute Error")
plt.legend()
plt.title("Training vs Validation Error")
plt.show()

```

```
# 10. Predict and Compare
# =====
y_pred = model.predict(x_test[:10]).flatten()
print("Predicted prices:", y_pred)
print("Actual prices:", y_test[:10].values)
```

Dataset shape: (20640, 9)

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	

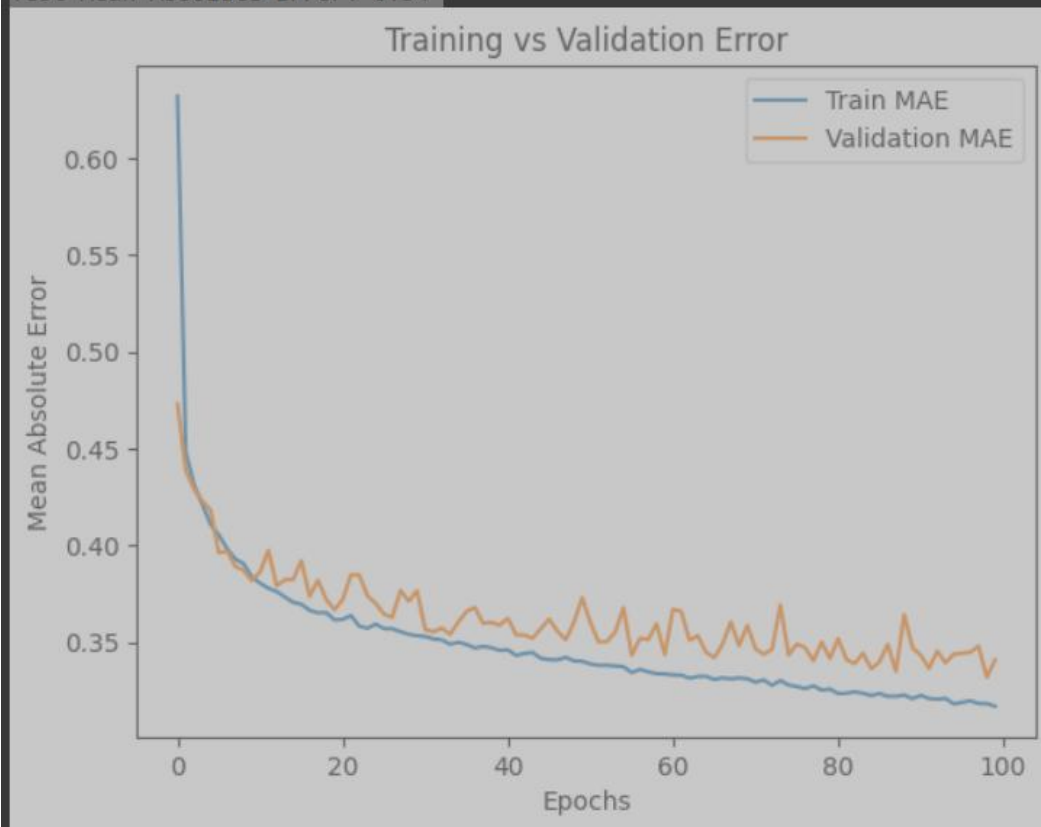
	Longitude	MEDV
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422

Dataset shape: (20640, 9)

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	

	Longitude	MEDV
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422

3 -122.25 3.413
4 -122.25 3.422
Train size: (16512, 8) Test size: (4128, 8)
Test Mean Absolute Error: 0.34



1/1 0s 75ms/step
Predicted prices: [0.4555246 0.8818768 5.2186923 2.4325957 2.9516323 1.6725786 2.128673
1.4970808 2.2517397 4.4329834]